June 24, 2026

Editors
IEEE Access

Dear Editors,

We are resubmitting our manuscript (Access-2026-23347) entitled

"When Agreement Fails: Visual Anchoring and Multimodal Pseudo-Label Reliability in Classroom Behavior Recognition"

following the reviewers' constructive feedback. We thank the Associate Editor and both reviewers for their careful reading. We have addressed all seven comments from Reviewer 2 in full; Reviewer 1 raised no technical concerns. A detailed point-by-point response accompanies this submission.

In addition to the reviewer-driven revisions (restructured Abstract, new Pipeline Overview, formal visual-anchoring definition, merged tables, and language polish throughout), the revised manuscript incorporates two new experiments completed during the revision period:

1. **Ten-model cross-model validation** (expanded from five models in the original submission), spanning three model families (Qwen, LLaVA, Gemma) across 7B–35B parameter scales. The expanded analysis reveals a clear capability boundary on intent-dependent categories and identifies Qwen3.5-27B as the most balanced annotator on the hardest dataset.
2. **Quality-threshold experiment** using Qwen3.5-27B pseudo-labels on TeacherBehavior, which locates the critical annotation accuracy range at which pseudo-label fine-tuning begins to exceed the zero-shot baseline—a finding with direct practical implications for MLLM-assisted annotation pipelines.

The revised manuscript makes four contributions. First, it provides a reproducible audit protocol linking annotation quality, downstream CLIP adaptation, retention controls, distribution-shift checks, anchor-proxy analysis, auxiliary baselines, and a ten-model cross-model robustness check. Second, it demonstrates that pseudo-label usefulness is conditional on category-level visual anchoring. Third, it shows that agreement filtering is not a portable denoising rule. Fourth, it identifies a practical annotation quality threshold and shows that the selective-routing gain scales with annotator quality.

We believe the revised manuscript is now substantially stronger and well suited to *IEEE Access*. All data are publicly available; the code repository (`https://github.com/zhanglizhuo/llm-annotation`) includes environment specifications, reproduction scripts, and tracked result files.

This submission is original, has not been published previously, and is not under consideration elsewhere. All authors have approved the revised version and agree to its resubmission to *IEEE Access*. The authors declare no competing interests.

Funding statement: This research was funded by the Education Department of Hunan Province, grant number HNJG-2022-0608.

Data and code availability: The SCB dataset is publicly available at `https://huggingface.co/datasets/wintonYF/SCB-Dataset`. The experiment code and results are available at `https://github.com/zhanglizhuo/llm-annotation`.

Ethics and privacy: This study is a secondary analysis of a publicly released benchmark. The manuscript reports aggregate results, makes no re-identification attempts, and uses privacy-preserving pixelated representative crops where visual examples are needed.

Thank you for your time and consideration.

Sincerely,

Yan Ma and Lizhuo Zhang

Corresponding author: Lizhuo Zhang
School of Information and Intelligent Science and Technology,
Hunan Agricultural University, Changsha 410128, China
Email: zhanglizhuo@hunau.edu.cn